

HW3

```
In [3]: import numpy as np
import pandas as pd
```

```
In [4]: data = pd.read_csv('titanic_data.csv')

y = data.Survived
X = data.loc[:, 'Pclass': "Fare"]
data.head()
```

Out[4]:

	Survived	Pclass	Sex	Age	Siblings/Spouses Aboard	Parents/Children Aboard	Fare
0	0	3	0	22.0	1	0	7.2500
1	1	1	1	38.0	1	0	71.2833
2	1	3	1	26.0	0	0	7.9250
3	1	1	1	35.0	1	0	53.1000
4	0	3	0	35.0	0	0	8.0500

```
In [5]: toy = data.head()
toy
```

Out[5]:

	Survived	Pclass	Sex	Age	Siblings/Spouses Aboard	Parents/Children Aboard	Fare
0	0	3	0	22.0	1	0	7.2500
1	1	1	1	38.0	1	0	71.2833
2	1	3	1	26.0	0	0	7.9250
3	1	1	1	35.0	1	0	53.1000
4	0	3	0	35.0	0	0	8.0500

3.1

Consider log-likelihood function:

$$\ell(\theta) = \sum_{i=1}^N y_i \log\left(\frac{1}{1 + e^{-\theta^\top x_i}}\right) + (1 - y_i) \log\left(\frac{1}{1 + e^{\theta^\top x_i}}\right)$$

```
In [6]: def LogLik(theta):
        return X.dot(theta).dot(y) - np.log(np.exp(1 + X.dot(theta))).sum()
        def Gr(theta):
            return (y - 1/(1 + np.exp(-X.dot(theta))))*X
```

```
In [7]: def GD(steps, size, tol):
        i = 0
        theta = np.linalg.inv(X.T.dot(X)).dot(X.T).dot(y)
        loglik = LogLik(theta)
        vec = [loglik]
        while i < steps:
            thetal = theta + size*Gr(theta)
            vec.append(LogLik(thetal))

            if np.square((thetal - theta)).sum() <= tol or np.square(vec[-1] - vec[-2])<= tol:
                print('Total %d' %i , 'iteration')
                return thetal, vec

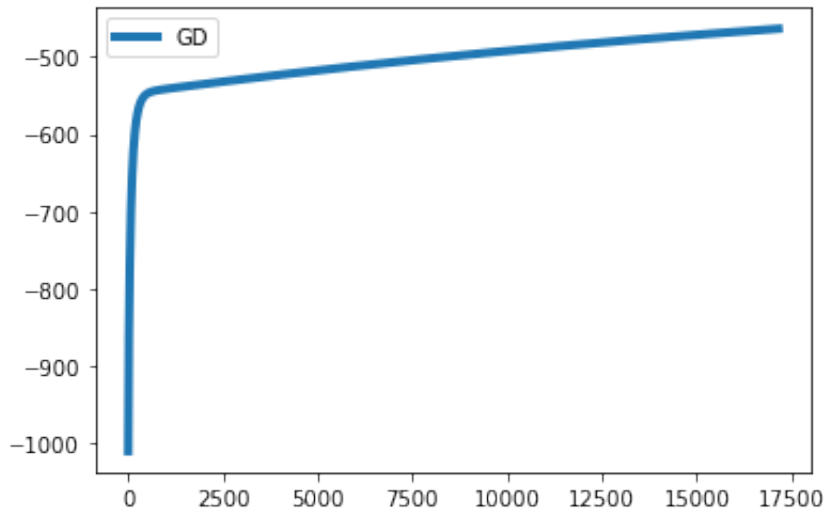
            theta = thetal
            i += 1

        print('Total %d' %i , 'iteration')
        return theta, vec
```

```
In [8]: theta, vec = GD(20000,1e-7, 1e-10)
```

Total 17155 iteration

```
In [9]: df = pd.DataFrame({"GD":vec})
df.plot.line(linewidth = 4)
None
```



```
In [10]: theta
```

```
Out[10]: Pclass      -0.157296
Sex           0.691840
Age          -0.022838
Siblings/Spouses Aboard -0.128911
Parents/Children Aboard -0.040576
Fare         0.013129
dtype: float64
```

(a)

I choose constant stepsize as 10^{-7} , the optimization process went pretty well. I've chosen some other strategies at the beginning, such as exact lone search/ learning rate decay, momentum method, or some other methods involving second gradient. The method involving learn search didn't get good result and usually requires more iterations. The likelihood will fluctuate in some iterations. The method methods involving second gradient cost more time and computation complexity.

In the end, the constant stepsize showed the desired result.

(b)

```
In [11]: import time
start_time = time.time()
theta, vec = GD(20000, 1e-7, 1e-10)
print("Total %s seconds elapsed" % (time.time() - start_time))
```

```
Total 17155 iteration
Total 28.299755096435547 seconds elapsed
```

It took 28.30 seconds to complete the iteration. It took about 3 seconds to converge according to the loglikelihood plot.

(c)

```
In [12]: theta
```

```
Out[12]: Pclass          -0.157296
Sex              0.691840
Age             -0.022838
Siblings/Spouses Aboard -0.128911
Parents/Children Aboard -0.040576
Fare            0.013129
dtype: float64
```

(d)

```
In [13]: vec[-1]
```

```
Out[13]: -463.53153422224244
```

(e)

Using θ^* to denote true parameter.

Consider Fisher information:

$$\mathbf{I}_{\theta^*} := -\mathbb{E} \left[\frac{d^2 \ell(\theta)}{d\theta^2} \Big|_{\theta=\theta^*} \right] = \mathbb{E} \left[\sum_{i=1}^N \frac{e^{-\theta^{\top} \mathbf{x}_i}}{(1 + e^{-\theta^{\top} \mathbf{x}_i})^2} \mathbf{x}_i \mathbf{x}_i^{\top} \Big|_{\theta=\theta^*} \right] = \sum_{i=1}^N \frac{e^{-\theta^{*\top} \mathbf{x}_i}}{(1 + e^{-\theta^{*\top} \mathbf{x}_i})^2} \mathbf{x}_i \mathbf{x}_i^{\top}$$

Let

$$\hat{\theta} := \arg \max_{\theta \in \mathbb{R}^{D+1}} \ell(\theta)$$

Suppose $\frac{d^2 \ell(\theta)}{d\theta^2}$. Then we have:

$$\hat{\theta} \xrightarrow{d} \mathcal{N}(\theta^*, \mathbf{I}_{\theta^*}^{-1})$$

According to computation, we have:

```
In [14]: def I(theta):
          W = np.diag(np.exp(-X.dot(theta))/(1 + np.exp(-X.dot(theta)))**2)
          return X.T.dot(W).dot(X)
```

```
In [15]: var = np.linalg.inv(I(theta))
```

Then we have:

$$\hat{\theta} \xrightarrow{d} \mathcal{N}(\theta^*, \mathbf{I}_{\hat{\theta}}^{-1})$$

where $\mathbf{I}_{\hat{\theta}}^{-1}$ is

In [16]: var

```
Out[16]: array([[ 3.21359720e-03, -1.96270908e-03, -1.94319179e-04,
                 -1.48733570e-03, -1.01101582e-03,  5.15191289e-05],
                [-1.96270908e-03,  2.37288881e-02, -3.72060099e-05,
                 1.08457571e-04, -2.95630615e-03, -7.20737802e-05],
                [-1.94319179e-04, -3.72060099e-05,  2.19491616e-05,
                 1.01031536e-04,  6.98258066e-05, -7.11340387e-06],
                [-1.48733570e-03,  1.08457571e-04,  1.01031536e-04,
                 5.81895119e-03, -2.23402533e-03, -6.34861528e-05],
                [-1.01101582e-03, -2.95630615e-03,  6.98258066e-05,
                 -2.23402533e-03,  1.05884034e-02, -5.48940130e-05],
                [ 5.15191289e-05, -7.20737802e-05, -7.11340387e-06,
                 -6.34861528e-05, -5.48940130e-05,  6.73377556e-06]])
```

In []:

3.2

Consider the the log-odds are defined as

$$\omega^* := \log\left(\frac{\mathbb{P}(y = 1 \mid \mathbf{x})}{\mathbb{P}(y = 0 \mid \mathbf{x})}\right) = \theta^{*\top} \mathbf{x}$$

By the invariance property of the MLE, we know that the MLE of ω^* is given by

$$\hat{\omega} := \hat{\theta}^\top \mathbf{x}$$

We further have:

$$\hat{\omega} \xrightarrow{d} \mathcal{N}(\omega^*, \mathbf{x}^\top \mathbf{I}_{\theta^*}^{-1} \mathbf{x})$$

In []:

3.3

My own feature would be Plass:3, Gender:0, Age:23, Siblings:0, Parents/Children:0, Fare: 7.25.

```
In [17]: x = np.array([3,0,23,0,0,7.25])
```

```
In [18]: ome_hat = 1 / (1 + np.exp(theta.dot(x)))
ome_hat
```

```
Out[18]: 0.7113562614023329
```

(a)

since $0.66 > 0.5$. According to (6.7) in lecture notes, I will survive the Titanic sinking.

(b)

According to:

$$\mathbb{P}(|\hat{\omega} - \omega^*| > \tau) = 2\Phi(\tau \mid 0, \mathbf{x}^\top \mathbf{I}_{\theta^*}^{-1} \mathbf{x})$$

```
In [19]: np.linalg.inv(I(theta)).dot(x).dot(x)
```

```
Out[19]: 0.013940140637597352
```

$$\tau = \Phi^{-1}(\alpha/2 \mid 0, \mathbf{x}^\top \mathbf{I}_{\theta^*}^{-1} \mathbf{x})$$

When $\alpha = 0.05$, the τ I compute based on my estimated variance is 0.2245956.

```
In [20]: ome_hat - 0.2245956
```

```
Out[20]: 0.4867606614023329
```

```
In [21]: ome_hat + 0.2245956
```

```
Out[21]: 0.9359518614023329
```

Then the 95% confidence interval would be [0.487,0.936]

(c)

I think my answer from (a) is fairly certain. In (a) I compute estimated probability larger than 1/2. In (c) I got my 95% confidence interval. Most of my interval lies in [0.5,1], so I think the interval I got coincides with the answer I got from (a).

In []:

3.4**(a)**

Consider LRT test:

$$\frac{\mathbb{P}(\hat{\theta}_j \mid \theta_j^* \neq 0)}{\mathbb{P}(\hat{\theta}_j \mid \theta_j^* = 0)} \geq_{H_0}^{H_1} 1$$

$$\left(\frac{\hat{\theta}_j}{v_j} \right)^2 \geq_{H_0}^{H_1} 2 \log \tau$$

The quantile I calculated for chi-square distribution with df 1 is 3.841459.

```
In [22]: np.square(theta/ np.diag(var)) > 3.841459
```

```
Out[22]: Pclass      True
         Sex         True
         Age         True
         Siblings/Spouses Aboard  True
         Parents/Children Aboard  True
         Fare         True
         dtype: bool
```


(b)

According the result in (a), all the features are significant.

(c)

Yes, the survival prediction is based on the feature prediction. If I change most significant feature in my feature vector, my survival prediction will change.

In []:

In []: