
COVID-19 Diagnosis based on CT Scans

AUTHORS: Yuchen Zeng, Zhuoyan Xu

We applied several machine learning methods on the scanned CT chest images to explore the relationship between positiveness of COVID-19 and chest CT.

We considered several competitive machine learning approaches, including Logistic Regression, K Nearest Neighbors, Naive Bayes, Decision Tree, Random Forest, Support Vector Machine and Convolutional Neural Network. We trained these models with raw image data and make predictions. We implemented the transfer learning on our own dataset. We used ResNet-18 architecture and personalized our own fully connection layer at the end. We froze the parameters in previous layers] only trained the personalized layers. In our result, we have SVM achieve the best performance with test accuracy 89%. The ResNet-18 has the second high test accuracy of 81.33%.

To combine the feature extraction advantage of convolutional neural networks and the interpretation of other machine learning methods, we used simple CNN to process images and transfer them into vectors representing the image characteristics. Then we applied several machine learning models to the image features. The advantage of this idea is instead of using raw image pixel value for classification, we implemented our machine learning algorithms on feature vectors extracted by pre-trained Convolutional Neural Network layers. This can reduce the feature dimension significantly, hence reduce the time and space complexity, free up more resources for more analysis. We aim to improve the accuracy of prediction and explore the explanation and interpretation as well.

We make use of the high accuracy from neural networks along with the interpretation from other machine learning models to better help make decisions and explain the relationship between COVID results and scanned images from a different perspective. We hope our model provides some insights into potential applications of machine learning methods in this specific COVID-19 situation.

1 Introduction

After China reported its first case of COVID-19 in December 2019, COVID-19 spread rapidly and became a global issue of public health since January 2020. According to the New York Times, 55.4 million people were tested positive for COVID-19, and 1.33 million people died as a result. COVID-19 is caused by the coronavirus SARS-CoV-2. The elderly and people with serious basic medical diseases (such as heart disease, lung disease or diabetes) seem to be more likely to have more serious complications from COVID-19. It is believed that COVID-19 is mainly spread through close contact between people, including people who are physically close to each other (within about 6 feet). People who are infected but have no symptoms can also spread the virus to other people.

As specific drugs for COVID-19 are not available, there is an urgent need to detect the disease early and isolate the infected person immediately. A single-slice of Computed tomography (CT) contains enough clinical information for accurate diagnosis decision-making. Bilateral changes in the chest CT scans can be observed from the clinical reports of infected people[1]. Due to its high sensitivity, chest CT has been used as an alternative tool to detect COVID-19 infection[2].

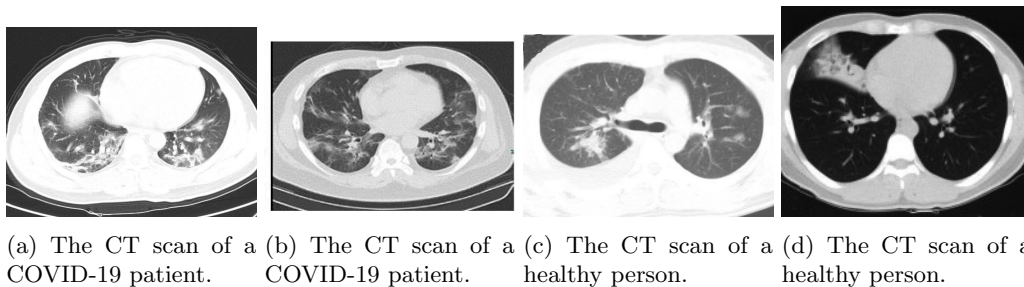


Figure 1: Examples of CT scans from COVID-CT dataset.

This work is to develop a classifier for accurate diagnosis of COVID-19 based on the CT scans. We train the classifier based on three popular machine learning methods: random forest, SVM and Neural network. Since neural network is the most commonly used method for image processing, we train two Neural network models: simple convolutional neural network and ResNet-18 to perform COVID-19 CT classification. Out of consideration of computation complexity, we use pre-trained ResNet-18 to improve the efficiency of ResNet-18. Finally, we perform numerical evaluations of the selected methods on the CT dataset, comparing and contrasting their performance of both accuracy and efficiency.

The highest accuracy we obtain is 89%, which is achieved by SVM classifier. The accuracy of our SVM classifier is higher than the baseline method *self-Tran+DenseNet-169* which is proposed by He et al.[3].

Contributions of this paper are summarized as follows. The rest of the paper is organized as follows. Section 2 discusses the related work on COVID-19 diagnosis based on CT scans. Section 3 introduce the dataset we analyze in details. We present the basic ideas of the selected methods Section 4, describe the experiment setting, report and discuss the empirical results in Section 5. We summarize the conclusions of this work, and discuss interesting future follow-up work in Section 6. The implementation details are provided in Appendix.

2 Related work

As coronavirus caused a global epidemic problem that spread quickly, there has been a fair amount of work in diagnosis of COVID-19[3, 4, 5, 6, 7, 1]. Diagnosis of COVID-19 is typically associated with symptoms of pneumonia, computed tomography (CT) scans[3, 7], and chest X-ray[4, 5, 6]. Our work is closely related to He et al.[3], Xu et al.[1] and Singh et al.[7] which also used CT scans to perform diagnosis of COVID-19. Specifically, we use the same dataset[8] as He et al.[3].

He et al.[3] combined many popular models in deep learning with transfer learning to train the data. He et al. proposed *Self-Trans*, a self-supervised transfer learning approach where contrastive self-supervised learning is integrated into transfer learning process to adjust the network weights pretrained on source data. The highest accuracy 86% is achieved by combining Self-Tran and DenseNet-160.

Xu et al.[1] evaluated two convolutional neural networks (CNN) three-dimensional classification models. One was ResNet-based network and another model was based on ResNet-based network structure by concatenating the location attention mechanism in the full-connection layer. It achieves an overall accuracy of 86.7%.

Singh et al.[7] used a CNN whose initial parameters of CNN are tuned using multiobjective differential evolution (MODE). The highest accuracy of the proposed approach achieves above 92%.

Ardakani et al.[9] implemented several well-known CNN architectures such as AlexNet, VGG, ResNet to

predict on CT scans. The best performance is achieved by ResNet101. They demonstrated in their study that the ResNet-101 can be considered as a promising model to characterize and diagnose COVID-19 infections.

In contrast, our work is not restricted to deep learning methods, but also consider other competitive machine learning approaches. We consider other benchmark machine learning methods: Random Forest and Support Vector Machine. We try to improved the accuracy of prediction and find explanation and interpretaion as well. We compare these methods in the Section 5.

3 Dataset

Data provided by Zhao et al.[8] has 746 CT images, containing clinical findings of COVID-19 from 216 patients. The images in the dataset consist of 349 CT scans that are positive for COVID-19 and 397 CT scans are negative for COVID-19. Figure 1 contains some sample images from the dataset. The size of images are different. The minimum width is 115 pixels, and the minimum height is 61 pixels.

4 Approach

In this section, we will describe which approaches we utilize to train our data and make classification.

4.1 Random forest

Random forests or random decision forests are one of the most widely used machine learning algorithms for regression and classification. Random forests are essentially a *bootstrap*, or *bagging* of *decision trees*.

Suppose the data matrix $\mathbf{X} = (\mathbf{x}_1^\top, \dots, \mathbf{x}_n^\top)^\top \in \mathbb{R}^{n \times d}$. A decision tree is a flowchart-like tree structure where the branch represents a decision rule, and each leaf node represents the outcome. A decision tree makes predictions based on series of questions. In practice, we train the decision tree by choosing the most informative feature via mutual information or other criterion in each internal node, and split the data according to this feature.

However, decision trees is sensitive to the data. Thus they can be quite biased, and tend to overfit. A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting. Random forests generally outperform decision trees.

4.2 Neural network

The idea of neural networks comes from the biological brains and neurons. Each neuron receives information from several previous neurons. In the human body, eyes and ears receive information, then the neurons transfer the information to the brain to make decisions. A node without an activation function is known as a perceptron.

One layer consists of multiple nodes. A standard neural network contains several layers: one input layer, one output layer, and several hidden layers. The input layer receives the feature vectors from each observation, while the output layer outputs the prediction probability of each class of label, and the hidden layer contains the intermediate results. The number of nodes in the input is the number of training samples in each batch, the number of nodes in the output layer is the number of labels, the number of nodes in hidden layers and

the number of hidden layers is determined by users based on the empirical performance. During the training period, the weights and bias are updated to minimize the selected objective function such as cross entropy via stochastic gradient descent.

A convolutional neural network adds several other layers especially for image processing. An image contains pixels represented by value. A convolutional layer can extract certain features in an image, such as eyes, mouth, ears of a dog. The word convolution comes from the functional analysis, it has an integral expression that can be used for computing the sum of two random variables. The convolution in CNN conducts the similar steps, it has several feature detectors, each detector walk through the image and compute the sum of the product of values stored in corresponding pixels.

We train the CT datasets with two CNN models: simple CNN and ResNet.

4.2.1 Simple CNN

We build one simple CNN model with three convolutional layers to train the data. To avoid the large parameter dominate the updates, we use the batch normalization in two convolutional layers. The max-pooling layers are introduced to help with local invariance. They introduce no parameters. The architecture of this model is shown as Figure 2.

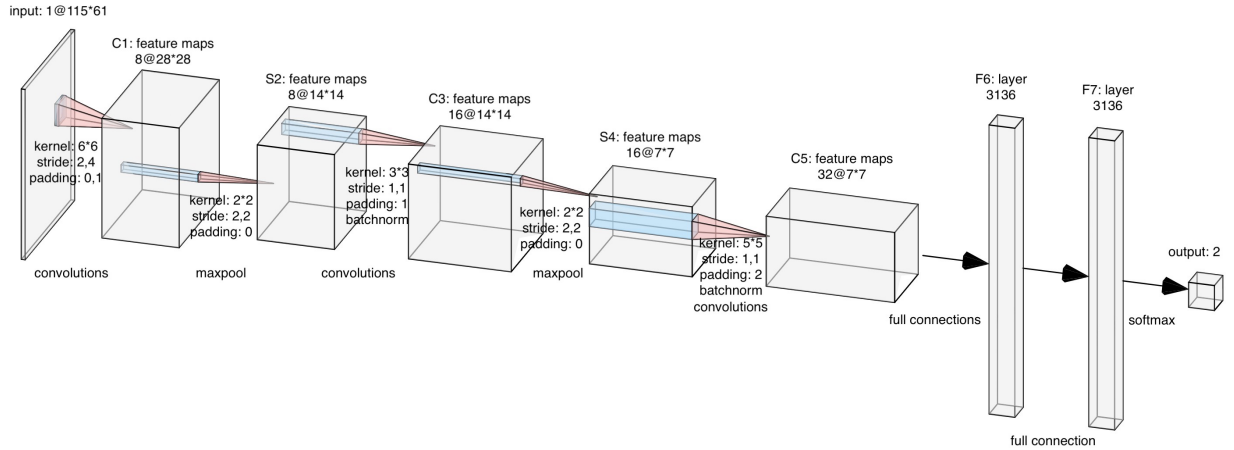


Figure 2: Architecture of the simple CNN we trained.

4.2.2 Transfer Learning + ResNets

The accuracy of the simple CNN we build is 70.67%, which is far less than the accuracy achieved by the baseline method He et al.[3]. We need a CNN model with a deeper architecture to improve the accuracy of classification. Deeper architecture significantly improves the accuracy over training set, with a cost of computation complexity. However, the heavy computation burden may make these model infeasible. We use transfer learning and ResNets to solve this issue.

ResNets is a CNN architecture proposed by He et al.[10]. The key idea of ResNets is that it allows to skip connections such as layers that are not useful. Therefore, ResNets allow us to implement significantly deep architectures. Transfer learning[11] is a technique proposed to mitigate the computation burden. The basic

idea is to store the knowledge gained in one training task and apply it to another. The knowledge in feature extraction layers of a well-trained model will help to extract features of an image on a new dataset. The features include the ears or mouths of a dog or cat in simple classification tasks. We stored the parameters in convolutional layers of a trained model and assigned them to our constructed models. We froze the parameters in those layers in training task and only update the last few fully connection layers.

4.3 Support vector machines (SVMs)

Support vector machines (SVMs) are considered as one of the best "out-of-box" classifier. Before SVM was proposed, perceptron is one of the most popular classification method. However, perceptron need high computation budget, and may have multiple solutions. The motivation of SVM is to solve the problems perceptron have when data is linearly separable.

Let $\mathbf{X} = (\mathbf{x}_1^\top, \dots, \mathbf{x}_n^\top)^\top \in \mathbb{R}^{n \times d}$ denote the sample matrix, where $\mathbf{x}_1, \dots, \mathbf{x}_n$ are individual samples. $\mathbf{y} = (y_1, \dots, y_n)$ are the labels, where $y_i \in \{0, 1\}$ for $i = 1, \dots, n$. The training of a SVM classifier involves finding a hyperplane

$$\mathcal{H}_{\beta, \beta_0} = \{i \in \{1, \dots, n\} : y_i(\langle \mathbf{x}_i, \beta \rangle + \beta_0) \geq 1\}, \quad \|\beta\| < m,$$

to separate the training samples from two different class with the largest margin m [12]. To maximize the margin m :

$$m = \min_{i=1, \dots, n} \text{dist}(\mathbf{x}_i, \mathcal{H}_{(\tilde{\beta}, \tilde{\beta}_0)}).$$

However, the requirement that the data is linear separable have been too strong, leading to many possible relaxations, e.g. soft-margin SVM. Soft-margin SVM state a preference for margins that classify the training data correctly, but soften the constraints to allow for non-separable data with a penalty $(\varepsilon_1, \dots, \varepsilon_n) \in \mathbb{R}^n$ proportional to the amount by which the example is misclassified. Thus the optimization problem becomes:

$$\begin{aligned} \min_{\beta, \beta_0, \varepsilon_1, \dots, \varepsilon_n} \quad & \frac{\|\beta\|^2}{2} + C \sum_{i=1}^n \varepsilon_i \\ \text{s.t.} \quad & y_i(\langle \mathbf{x}_i, \beta \rangle + \beta_0) \geq 1 - \varepsilon_i \quad i = 1, \dots, n \\ & \varepsilon_i \geq 0 \quad i = 1, \dots, n \end{aligned}$$

In practice, C is determined by cross validation.

A more advanced tool is kernel, which transforms the data to be linear separable. Suppose the transformation $\psi : \mathbb{R}^d \rightarrow \text{Hilbert space } H$ transforms the data $\mathbf{x}_1, \dots, \mathbf{x}_n$ to linear separable data $\psi(\mathbf{x}_1), \dots, \psi(\mathbf{x}_n)$, then the primal optimization problem becomes:

$$\begin{aligned} \min_{\beta \in H, \beta_0 \in \mathbb{R}} \quad & \frac{\|\beta\|_H^2}{2} + C \sum_{i=1}^n \varepsilon_i \\ \text{s.t.} \quad & y_i(\langle \psi(\mathbf{x}_i), \beta \rangle_H + \beta_0) \geq 1 - \varepsilon_i \quad i = 1, \dots, n \\ & \varepsilon_i \geq 0 \quad i = 1, \dots, n \end{aligned}$$

The dual problem is

$$\begin{aligned} \max_{\lambda \in \mathbb{R}^n} \quad & -\frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j \langle \psi(\mathbf{x}_i), \psi(\mathbf{x}_j) \rangle_H + \sum_{i=1}^n \lambda_i \\ \text{s.t.} \quad & \sum_{i=1}^n \lambda_i y_i = 0, \quad \lambda_i \geq 0 \quad i = 1, \dots, n. \end{aligned}$$

With kernel $K = \langle \psi(\cdot), \psi(\cdot) \rangle_H$, the dual problem can be represented as

$$\begin{aligned} \max_{\lambda \in \mathbb{R}^n} & -\frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j K(\mathbf{x}_i, \mathbf{x}_j) + \sum_{i=1}^n \lambda_i \\ \text{s.t.} & \sum_{i=1}^n \lambda_i y_i = 0, \lambda_i \geq 0 \quad i = 1, \dots, n. \end{aligned}$$

Actually, given K , ψ is unnecessary for us to solve the dual problem. In practice, we select the kernel K and then solve the dual problem to get SVM classifier.

4.4 Simple CNN as feature extractor

We also consider using simple CNN to process images and transfer them into vectors representing the image characteristics, then we use other machine learning methods in previous sections on the feature vectors. The main idea is similar to transfer learning, the convolutional layer is suitable for extracting features in images. Ideally, applying previous machine learning methods on these vectors can give higher accuracy and better interpretation than applying them on raw images.

We use the CNN architecture in 4.2.1 as feature extractor, we add one layer with 10 nodes right before the output layer. Once we finish the training, we feed all images to the architecture and get the intermediate result outputted by the second to last layer. Then we have the new dataset with each observation has one corresponding feature vector.

5 Experiments

5.1 Main result on raw images

We resize all images to size 115×61 , and divide them into training set and testing set at a ratio of 9:1. The packages we use are listed below:

- file reading: `pandas`, `skimage`, `shutil`, `os`;
- data processing: `pandas`, `numpy`, `random`, `sklearn`, `PIL`;
- benchmarking: `time`;
- plot: `matplotlib`;
- classifier: `sklearn`, `torch`, `torchvision`.

We provide more details of how we use these packages in Appendix. We train the methods we introduced in last section, and the settings are described below.

Random Forest: We flatten the image arrays to vectors of length 7015. Our random forests generate 100 decision trees, whose sample size is the same as the original sample size but the samples are drawn with replacement. Our criteria of choosing features in each internal node is Gini impurity. Gini impurity is also known as the total decrease in node impurity. This is how much the model fit or accuracy decreases when you drop a variable. The test accuracy of our random forest is around 80%.

Simple CNN: The learning rate, batch size and number of epochs we select empirically are 0.05, 8 and 20, respectively. The objective function we use here is cross entropy. All the activation functions except the output layer are ReLU. The test accuracy we obtain from simple CNN is 70.67%.

Transfer Learning+ResNet: We use a pre-trained ResNet-18[13] and only train the last hidden layer with activation function ReLU and output layer. To avoid overfitting, we drop feature maps with a probability of 0.5. The accuracy of the training set and test set are 82.116% and 81.33%, respectively, indicating that there is almost no overfitting.

Support Vector Machine: We flatten the image matrices to vectors of length $115 \times 61 = 7015$. We use 3-fold cross-validated grid search to determine the kernel from two alternative kernels, radial basis function kernel $\exp(-\gamma\|\mathbf{x} - \mathbf{x}'\|^2)$ and linear kernel $\mathbf{x}^\top \mathbf{x}'$, the $\gamma \in \{0.001, 0.0001\}$ in radial basis function kernel and the regularization coefficient $C \in \{1, 10, 100, 1000\}$. Therefore, we choose radial basis function kernel with $\gamma = 0.001$, $C = 10$. The hyperparameters we obtain are $C = 10$, $\gamma = 0.001$, and selected kernel is radial basis function kernel. The test accuracy our SVM classifier we obtain is 89%.

We summarize the empirical performance of the selected approaches as Table 2. Both SVM and Random Forest achieves 100% accuracy on training set. However, the test accuracy of random forest is only 77%, which is far less than 100%, thus we can conclude that the random forest suffers from serious over-fitting. The highest accuracy is achieved by SVM, which is 89%, higher than the baseline method He et al.[3]. SVM also stand out for its computation efficiency. The execution time of SVM is only 0.87 minutes, far less than another competitive model: ResNet-18, which takes 12.58 minutes to train the model.

Model	Train Accuracy	Test Accuracy	Execution time (min)
SVM	100%	89%	0.87
Random Forest	100%	77%	0.04
Simple CNN	81.222%	70.67%	4.61
ResNet-18	82.116%	81.33%	12.58
He et al.[3]	-	86%	-

Table 1: Comparison of SVM, random forest, simple CNN, ResNet-18 and the baseline method He et al.[3].

5.2 Results on extracted feature by simple CNN

We use the idea in 4.4, we apply several machine learning methods on the extracted feature vectors, we also search for the best parameters for each model and make predictions. We divided the datasets to 5 fold and compute the mean cross validation accuracy. The results are shown below:

Model	Train Accuracy	Test Accuracy	K-fold accuracy
Logistic Regression	99%	59%	97.52%
Decision Tree	98.14%	60%	97.37%
Random Forest	98.45%	57%	97.83%
K Nearest Neighbors	98.61%	58%	98.45%
Naive Bayes	98.60%	57%	98.45%
SVM	98.45%	63%	98.76%

Table 2: Comparison of the baseline method on extracted features.

The SVM gave the best test accuracy so we draw the hyperplane on the 2 dimension plot to show the region of data in 2 classes:

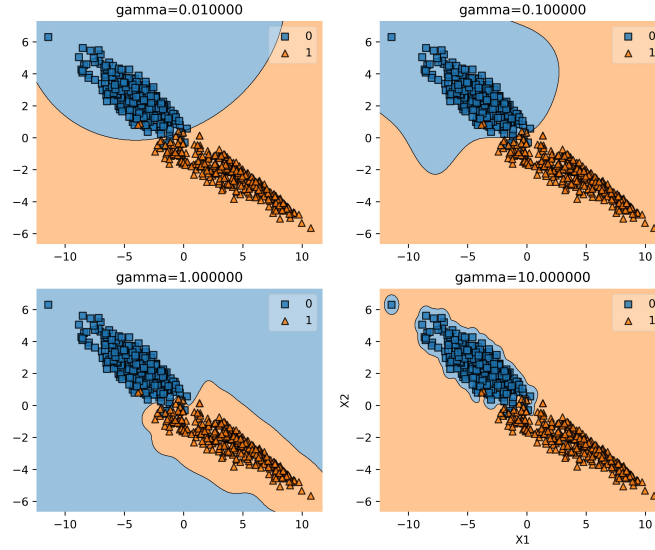


Figure 3: SVM with different γ in Gaussian kernel

The plot shows a fairly clear hyperplane to separate the two regions. We conclude the SVM can separate the data well on a higher dimensional space.

6 Conclusions and future work

We utilize the high accuracy from neural network along with the interpretation from machine learning models to better help make decisions and explain the relationship between COVID results and scanned images from a different perspective.

To summarize, SVM is the best model among the selected models to diagnose COVID-19, with the highest test accuracy of 89% and high computation efficiency. The ResNet architecture has the second high test accuracy, which can be improved if we tuned the layers and nodes to make them more suitable for the data.

The model on raw images works pretty well on the test dataset. However, the models on feature vectors have high train accuracy but lower test accuracy. The main reason is the potential overfitting problem when we fed the extracted features to our model. The CNN extract the main features but neglect some other important factors of the CT images. The focus on minor features may lose the overall characteristics of the images, we need to be careful when we use this idea in image classification.

In this work, the methods with raw dataset have generally better performance than those on extracted features. In our future work, we can explore more on how to use CNN architecture extract features well and combine them with our other machine learning models.

We will also focus on the feature importance generated by machine learning methods. We will explore which feature can best represent the CT image and help distinguish the Covid and Non-Covid results.

References

- [1] Xiaowei Xu, Xiangao Jiang, Chunlian Ma, Peng Du, Xukun Li, Shuangzhi Lv, Liang Yu, Yanfei Chen, Junwei Su, Guanqing Lang, Yongtao Li, Hong Zhao, Kaijin Xu, Lingxiang Ruan, and Wei Wu. Deep learning system to screen coronavirus disease 2019 pneumonia, 2020.
- [2] Xingzhi Xie, Zheng Zhong, Wei Zhao, Chao Zheng, Fei Wang, and Jun Liu. Chest ct for typical coronavirus disease 2019 (covid-19) pneumonia: Relationship to negative rt-pcr testing. *Radiology*, 2020.
- [3] Xuehai He, Xingyi Yang, Shanghang Zhang, Jinyu Zhao, Yichen Zhang, Eric Xing, and Pengtao Xie. Sample-efficient deep learning for covid-19 diagnosis based on ct scans. *medRxiv*, 2020.
- [4] Şaban Öztürk, Umut Özkaya, and Mücahid Barstuğan. Classification of coronavirus (covid-19) from x-ray and ct images using shrunken features. *International Journal of Imaging Systems and Technology*.
- [5] Terry Gao. Chest x-ray image analysis and classification for covid-19 pneumonia detection using deep cnn. *medRxiv*, 2020.
- [6] Asmaa Abbas, Mohammed Abdelsamea, and Mohamed Gaber. Classification of covid-19 in chest x-ray images using detrac deep convolutional neural network. *medRxiv*, 2020.
- [7] Dilbag Singh, Vijay Kumar, Vaishali, and Manjit Kaur. Classification of covid-19 patients from chest ct images using multi-objective differential evolution-based convolutional neural networks. *European Journal of Clinical Microbiology & Infectious Diseases*, 39(7):1379–1389, 2020.
- [8] Jinyu Zhao, Yichen Zhang, Xuehai He, and Pengtao Xie. Covid-ct-dataset: a ct scan dataset about covid-19. *arXiv preprint arXiv:2003.13865*, 2020.
- [9] Ali Abbasian Ardakani, Alireza Rajabzadeh Kanafi, U Rajendra Acharya, Nazanin Khadem, and Afshin Mohammadi. Application of deep learning technique to manage covid-19 in routine clinical practice using ct images: Results of 10 convolutional neural networks. *Computers in Biology and Medicine*, page 103795, 2020.
- [10] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016.
- [11] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition, 2015.
- [12] Vladimir N. Vapnik. *The Nature of Statistical Learning Theory*. Springer-Verlag, Berlin, Heidelberg, 1995.
- [13] Mujadded Al Rabbani Alif, Sabbir Ahmed, and Muhammad Hasan. Isolated bangla handwritten character recognition with convolutional neural network. pages 1–6, 12 2017.

A Appendix: Data processing Python source code

```

1 # pandas: plot the accuracy results, read files, save the data, etc
2 import pandas as pd
3
4 # os: make sure the code can be executed for all operating systems
5 import os
6
7 # random: draw random samples

```

```

8 import random
9
10 # imread: read and display images
11 from skimage.io import imread
12
13 # train_test_split: create validation set
14 from sklearn.model_selection import train_test_split
15
16 # pyplot: plot
17 import matplotlib.pyplot as plt
18
19 # Image: process images such as resize/clip the images
20 from PIL import Image
21
22 # shutil: remove an entire directory
23 import shutil
24
25 # load the images' names and label, and save them into a dataframe
26 # use filter to skip the files such as DS_Store and other caches
27 covidImgs = list(filter(lambda x: len(x.split(".")) > 1, os.listdir(os.path.join("dataset",
28 "CT", "CT_COVID"))))
29 healthImgs = list(filter(lambda x: len(x.split(".")) > 1, os.listdir(os.path.join("dataset",
30 "CT", "CT_NonCOVID"))))
31 # 1: True; 0: False
32 ctLabel = pd.DataFrame({"Image": covidImgs + healthImgs, "Covid": [1]*len(covidImgs) + [0]*
33 len(healthImgs)})
34 ctLabel
35
36 # create the training set and test set
37 train_name, test_name, train_y, test_y = train_test_split(ctLabel.Image, ctLabel.Covid,
38 test_size = 0.1)
39
40 # create the validation set
41 train_name, val_name, train_y, val_y = train_test_split(train_name, train_y, test_size =
42 0.1)
43
44 def resizeImage(folder):
45
46     """
47     Resize all the images in the folder to (115, 61).
48
49     Parameter
50     -----
51     folder: the folder to save the resized images: "train", "test" and "validation".
52
53     Return
54     -----
55     a list of image arrays
56     """
57
58     if folder == 'train':
59         x, y = train_name, train_y
60     elif folder == 'test':
61         x, y = test_name, test_y
62     else:
63         x, y = val_name, val_y
64
65     imgs = []
66     new_dir = os.path.join("dataset", "CT", folder)
67
68     # remove previous data, recreate the train/val/test dataset
69     if os.path.exists(new_dir) and os.path.isdir(new_dir):
70         shutil.rmtree(new_dir)
71         os.makedirs(new_dir)

```

```

67     for name, covid in zip(x, y):
68         original_folder = "CT_COVID" if covid else "CT_NonCOVID"
69         img_path = os.path.join("dataset", "CT", original_folder, name)
70         img = Image.open(img_path).convert('RGB').resize((115, 61)).convert('L')
71         new_path = os.path.join(new_dir, name)
72         img.save(new_path, "JPEG", optimize=True)
73         img.close()
74
75         #####
76         ##### read the image arrays (not necessary for this function) #####
77         ##### can be removed #####
78         #####
79         img = imread(new_path, as_gray=True)
80         # convert the type of pixel to float32
81         img = img.astype('float32')
82         # normalize the pixel values
83         img /= 255.0
84         # append the image into the list
85         imgs.append(img)
86
87     return imgs
88
89     train_x, test_x, val_x = resizeImage("train"), resizeImage("test"), resizeImage("
validation")
90
91     # display the CT images
92     for i in random.choices(range(len(train_x)), k=4):
93         plt.imshow(train_x[i], cmap='gray')
94         plt.show()
95
96     # write the data into csv files
97     pd.DataFrame({"name":train_name , "covid":train_y}).to_csv('train.csv', index=False)
98
99     pd.DataFrame({"name":test_name , "covid":test_y}).to_csv('test.csv', index=False)
100
101     pd.DataFrame({"name":val_name , "covid":val_y}).to_csv('validation.csv', index=False)

```

B Appendix: Random forest implementation with Python

```

1  # os: make sure the code related to path can work for all operating systems
2  import os
3
4  # time: benchmarking
5  import time
6
7  # sklearn: report the classification accuracy
8  import sklearn
9
10 # RandomForestClassifier: the random forest classifier
11 from sklearn.ensemble import RandomForestClassifier
12
13 # GridSearchCV: select the hyperparameter
14 from sklearn.model_selection import GridSearchCV
15
16 # pandas: plot the accuracy results, read files, save the data, etc
17 import pandas as pd
18
19 # numpy: process arrays
20 import numpy as np
21
22 # os: make sure the code can be executed for all operating systems

```

```

23 import os
24
25 # random: draw random samples
26 import random
27
28 # imread: read and display images
29 from skimage.io import imread
30
31 # train_test_split: create validation set
32 from sklearn.model_selection import train_test_split
33
34 # pyplot: plot
35 import matplotlib.pyplot as plt
36
37 # Image: process images such as resize/clip the images
38 from PIL import Image
39
40 # shutil: remove an entire directory
41 import shutil
42
43 # time: benchmarking
44 import time
45
46 DIMENSION = (115, 61)
47 IMG_PATH = os.path.join("dataset", "CT")
48
49 # Importing Train and Test datasets
50 train_data = pd.read_csv("train.csv")
51 test_data = pd.read_csv("test.csv")
52 val_data = pd.read_csv("validation.csv")
53
54 def readImage(folder):
55     """
56     Get the image data array of all images in one folder.
57
58     Parameters
59     -----
60     folder: the name of the folder: "train", "test", "validation".
61
62     Returns
63     -----
64     A list of image arrays and the labels.
65     """
66
67     data = pd.read_csv(folder + ".csv")
68     x, y = data.name, data.covid
69
70     imgs = []
71     new_dir = os.path.join(IMG_PATH, folder)
72
73     for name, covid in zip(x, y):
74         img_path = os.path.join(new_dir, name)
75         img = imread(img_path, as_gray=True).flatten()
76         # convert the type of pixel to float32
77         img = img.astype('float32')
78         # normalize the pixel values
79         img /= 255.0
80         # append the image into the list
81         imgs.append(img)
82
83     return imgs, y
84
85 # load the datasets
86 train_x, train_y = readImage("train")

```

```

87 test_x, test_y = readImage("test")
88 val_x, val_y = readImage("validation")
89
90 # ===== Using Random Forest without hyper paramter tuning and clustering
91 # =====
92 start = time.time()
93 rf = RandomForestClassifier(n_estimators = 100)
94 rf.fit(train_x+val_x, np.concatenate((train_y.values, val_y.values)))
95 # rf.fit(train_x, train_y)
96
97 print("Training data metrics:")
98 print(sklearn.metrics.classification_report(y_true = train_y, y_pred = rf.predict(train_x)))
99
100 print("Validation data metrics:")
101 print(sklearn.metrics.classification_report(y_true = val_y, y_pred = rf.predict(val_x)))
102
103 # Predictions on testset
104 # test data metrics
105 print("Test data metrics:")
106 print(sklearn.metrics.classification_report(y_true = test_y, y_pred = rf.predict(test_x)))
107 end = time.time()
108 print("Time elapsed: %.2f min" % ((end-start)/60))

```

C Appendix: Simple CNN implementation with Python

```

1 # In[1]:
2
3
4 # pandas: plot the accuracy results, read files, save the data, etc
5 import pandas as pd
6
7 # os: make sure the code can be executed for all operating systems
8 import os
9
10 # random: draw random samples
11 import random
12
13 # imread: read and display images
14 from skimage.io import imread
15
16 # train_test_split: create validation set
17 from sklearn.model_selection import train_test_split
18
19 # pyplot: plot
20 import matplotlib.pyplot as plt
21
22 # Image: process images such as resize/clip the images
23 from PIL import Image
24
25 # shutil: remove an entire directory
26 import shutil
27
28 # time: benchmarking
29 import time
30
31 # pytorch libraries and modules
32 import torch
33
34 # numpy: process the data
35 import numpy as np
36

```

```

37 from torch.autograd import Variable
38 from torch.nn import Linear, ReLU, CrossEntropyLoss, Sequential, Conv2d, MaxPool2d, Module,
    Softmax, BatchNorm2d, Dropout
39 from torch.utils.data import Dataset
40 from torchvision import transforms
41 from torch.utils.data import DataLoader
42 import torch.nn.functional as F
43 from torch.optim import Adam, SGD
44
45
46 # In[2]:
47
48
49 #-----
50 ### SETTINGS
51 #-----
52
53 # Hyperparameters
54 RANDOM_SEED = 1
55 LEARNING_RATE = 0.05
56 BATCH_SIZE = 8
57 NUM_EPOCHS = 20
58
59 # Architecture
60 NUM_CLASSES = 2
61 DIMENSION = (115, 61)
62
63 # Other
64 DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
65 IMG_PATH = os.path.join("dataset", "CT")
66 NUM_WORKERS = 0
67
68
69 # In[3]:
70
71
72 def readImage(folder):
73     """
74     Get the image data array of all images in one folder.
75
76     Parameters
77     -----
78     folder: the name of the folder: "train", "test", "validation".
79
80     Returns
81     -----
82     A list of image arrays and the labels.
83     """
84
85     data = pd.read_csv(folder + ".csv")
86     x, y = data.name, data.covid
87
88     imgs = []
89     new_dir = os.path.join(IMG_PATH, folder)
90
91     for name, covid in zip(x, y):
92         img_path = os.path.join(new_dir, name)
93         img = imread(img_path, as_gray=True)
94         # convert the type of pixel to float32
95         img = img.astype('float32')
96         # normalize the pixel values
97         img /= 255.0
98         # append the image into the list
99         imgs.append(img)

```

```

100
101     return imgs, y
102
103 # load the datasets
104 train_x, train_y = readImage("train")
105 test_x, test_y = readImage("test")
106 val_x, val_y = readImage("validation")
107
108
109 # In[4]:
110
111
112 plt.imshow(train_x[random.randrange(len(train_x))], cmap='gray')
113 plt.show()
114
115
116 # In[5]:
117
118
119 class CTDataset(Dataset):
120     """Custom Dataset for loading CT images"""
121
122     def __init__(self, csv_path, img_dir, transform=None):
123
124         df = pd.read_csv(csv_path)
125         self.img_dir = [img_dir] * len(df['name'].values)
126         self.img_names = df['name'].values
127         self.y = df['covid'].values
128         self.transform = transform
129
130     def __getitem__(self, index):
131         img = Image.open(os.path.join(self.img_dir[index],
132                                       self.img_names[index]))
133         # img = imread(os.path.join(self.img_dir[index],
134                                   self.img_names[index]), as_gray=True).astype('
135 float32')
136         if self.transform is not None:
137             img = self.transform(img)
138
139         label = self.y[index]
140         return img, label
141
142     def __len__(self):
143         return self.y.shape[0]
144
145     # concatenate two datasets
146     def __add__(self, newDataset):
147         self.img_names = np.concatenate((self.img_names, newDataset.img_names))
148         self.img_dir += newDataset.img_dir
149         self.y = np.concatenate((self.y, newDataset.y))
150         return self
151
152 # In[6]:
153
154
155 # Note that transforms.ToTensor()
156 # already divides pixels by 255. internally
157
158 custom_transform = transforms.Compose([#transforms.Lambda(lambda x: x/255.),
159                                       transforms.ToTensor()])
160
161
162 train_dataset = CTDataset(csv_path='train.csv',

```

```

163         img_dir=os.path.join("dataset", "CT", "train"),
164         transform=custom_transform)
165
166 valid_dataset = CTDataset(csv_path='validation.csv',
167                             img_dir=os.path.join("dataset", "CT", "validation"),
168                             transform=custom_transform)
169
170 # this method don't need validation dataset
171
172 train_loader = DataLoader(dataset= train_dataset+valid_dataset,
173                             batch_size=BATCH_SIZE,
174                             shuffle=True,
175                             num_workers=NUM_WORKERS)
176
177
178 valid_loader = DataLoader(dataset=valid_dataset,
179                             batch_size=BATCH_SIZE,
180                             shuffle=False,
181                             num_workers=NUM_WORKERS)
182
183 test_dataset = CTDataset(csv_path='test.csv',
184                             img_dir=os.path.join("dataset", "CT", "test"),
185                             transform=custom_transform)
186
187 test_loader = DataLoader(dataset=test_dataset,
188                             batch_size=BATCH_SIZE,
189                             shuffle=False,
190                             num_workers=NUM_WORKERS)
191
192
193 # In[7]:
194
195
196 # Checking the dataset
197 for images, labels in train_loader:
198     print('Image batch dimensions:', images.shape)
199     print('Image label dimensions:', labels.shape)
200     break
201
202 # Checking the dataset
203 for images, labels in train_loader:
204     print('Image batch dimensions:', images.shape)
205     print('Image label dimensions:', labels.shape)
206     break
207
208
209 # In[8]:
210
211
212 # This cell just checks if the dataset can be loaded correctly.
213
214 torch.manual_seed(0)
215
216 num_epochs = 2
217 for epoch in range(num_epochs):
218
219     for batch_idx, (x, y) in enumerate(train_loader):
220
221         print('Epoch:', epoch+1, end='')
222         print(' | Batch index:', batch_idx, end='')
223         print(' | Batch size:', y.size()[0])
224
225         x = x.to(DEVICE)
226         y = y.to(DEVICE)

```



```

227         print('break minibatch for-loop')
228         break
229
230
231
232 # # Multilayer Perceptron Model
233
234 # In[9]:
235
236
237 #####
238 ### NO NEED TO CHANGE THIS CELL
239 #####
240
241 def compute_epoch_loss(model, data_loader):
242     model.eval()
243     curr_loss, num_examples = 0., 0
244     with torch.no_grad():
245         for features, targets in data_loader:
246             features = features.to(DEVICE)
247             targets = targets.to(DEVICE)
248             logits, probas = model(features)
249             loss = F.cross_entropy(logits, targets, reduction='sum')
250             num_examples += targets.size(0)
251             curr_loss += loss
252
253     curr_loss = curr_loss / num_examples
254     return curr_loss
255
256
257 def compute_accuracy(model, data_loader, device):
258     model.eval()
259     correct_pred, num_examples = 0, 0
260     for i, (features, targets) in enumerate(data_loader):
261
262         features = features.to(device)
263         targets = targets.to(device)
264
265         logits, probas = model(features)
266         _, predicted_labels = torch.max(probas, 1)
267         num_examples += targets.size(0)
268         correct_pred += (predicted_labels == targets).sum()
269     return correct_pred.float()/num_examples * 100
270
271
272 # In[10]:
273
274
275 class ConvNet(torch.nn.Module):
276
277     def __init__(self, num_classes):
278         super(ConvNet, self).__init__()
279
280         self.num_classes = num_classes
281
282         ### Layers: ADD ADDITIONAL LAYERS BELOW IF YOU LIKE
283
284         # 115*61*1 => 28*28*8
285         self.conv_1 = torch.nn.Conv2d(in_channels=1, out_channels=8, kernel_size=(6,6),
stride=(2,4), padding=(0,1))
286
287         # 28*28*8 => 14*14*8
288         self.pool_1 = torch.nn.MaxPool2d(kernel_size=(2,2), stride=(2,2), padding=0)
289

```

```

290     # 14*14*8 => 14*14*16
291     self.conv_2 = torch.nn.Conv2d(in_channels=8, out_channels=16, kernel_size=(3,3),
stride=(1,1), padding=1)
292
293     # 14*14*16 => 7*7*16
294     self.pool_2 = torch.nn.MaxPool2d(kernel_size=(2,2), stride=(2,2), padding=0)
295
296     self.bn_2 = torch.nn.BatchNorm2d(16)
297
298     # 7*7*16 => 7*7*32
299     self.conv_3 = torch.nn.Conv2d(in_channels=16, out_channels=32, kernel_size=5, stride
=1, padding=2)
300
301     self.bn_3 = torch.nn.BatchNorm2d(32)
302
303     # Multilayer perceptron
304     self.linear_1 = torch.nn.Linear(7*7*32, 7*7*64)
305     self.bn_l1 = torch.nn.BatchNorm1d(7*7*64)
306     self.linear_2 = torch.nn.Linear(7*7*64, 7*7*64)
307     self.bn_l2 = torch.nn.BatchNorm1d(7*7*64)
308     self.linear_out = torch.nn.Linear(7*7*64, num_classes)
309
310
311     def forward(self, x):
312
313         ### MAKE SURE YOU CONNECT THE LAYERS PROPERLY IF YOU CHANGED
314         ### ANYTHNG IN THE __init__ METHOD ABOVE
315         out = self.conv_1(x)
316         out = F.relu(out)
317         out = self.pool_1(out)
318
319         out = self.conv_2(out)
320         out = self.bn_2(out)
321         # out = F.dropout(out, p=0.2, training=self.training)
322         out = F.relu(out)
323         out = self.pool_2(out)
324
325         out = self.conv_3(out)
326         out = self.bn_3(out)
327         # out = F.dropout(out, p=0.2, training=self.training)
328         out = F.relu(out)
329
330         out = self.linear_1(out.view(-1, 7*7*32))
331         out = self.bn_l1(out)
332         out = F.relu(out)
333         # out = F.dropout(out, p=0.2, training=self.training)
334
335         out = self.linear_2(out)
336         out = self.bn_l2(out)
337         out = F.relu(out)
338         # out = F.dropout(out, p=0.2, training=self.training)
339
340         logits = self.linear_out(out)
341         probas = F.softmax(logits, dim=1)
342         return logits, probas
343
344
345     #####
346     ### Model Initialization      ###
347     #####
348
349
350     # the random seed makes sure that the random weight initialization
351     # in the model is always the same.

```

```

352 # In practice, some weights don't work well, and we may also want
353 # to try different random seeds. In this homework, this is not
354 # necessary.
355 torch.manual_seed(RANDOM_SEED)
356
357 ### IF YOU CHANGED THE ARCHITECTURE ABOVE, MAKE SURE YOU
358 ### ACCOUNT FOR IT VIA THE PARAMETERS BELOW. I.e., if you
359 ### added a second hidden layer, you may want to add a
360 ### hidden_2 parameter here. Also you may want to play
361 ### with the number of hidden units.
362
363 model = ConvNet(NUM_CLASSES)
364 model = model.to(DEVICE)
365
366
367 ### For this homework, do not change the optimizer. However, you
368 ### likely want to experiment with the learning rate!
369 optimizer = torch.optim.SGD(model.parameters(), lr=LEARNING_RATE)
370
371
372 # In[11]:
373
374
375 #####
376 ### NO NEED TO CHANGE THIS CELL
377 #####
378
379 def train(model, train_loader, test_loader):
380
381     minibatch_cost, epoch_cost = [], []
382     start_time = time.time()
383     for epoch in range(NUM_EPOCHS):
384
385         model.train()
386         for batch_idx, (features, targets) in enumerate(train_loader):
387
388             features = features.to(DEVICE)
389             targets = targets.to(DEVICE)
390
391             ### FORWARD AND BACK PROP
392             logits, probas = model(features)
393             cost = F.cross_entropy(logits, targets)
394             optimizer.zero_grad()
395
396             cost.backward()
397             minibatch_cost.append(cost)
398
399             ### UPDATE MODEL PARAMETERS
400             optimizer.step()
401
402             ### LOGGING
403             if not batch_idx % 150:
404                 print ('Epoch: %03d/%03d | Batch %04d/%04d | Cost: %.4f'
405                       %(epoch+1, NUM_EPOCHS, batch_idx,
406                         len(train_loader), cost))
407
408
409         with torch.set_grad_enabled(False): # save memory during inference
410             print('Epoch: %03d/%03d | Train: %.3f%%' % (
411                 epoch+1, NUM_EPOCHS,
412                 compute_accuracy(model, train_loader, device=DEVICE)))
413
414             cost = compute_epoch_loss(model, train_loader)
415             epoch_cost.append(cost)

```

```

416         print('Time elapsed: %.2f min' % ((time.time() - start_time)/60))
417
418     print('Total Training Time: %.2f min' % ((time.time() - start_time)/60))
419
420
421
422     with torch.set_grad_enabled(False): # save memory during inference
423         print('Test accuracy: %.2f%%' % (compute_accuracy(model, test_loader, device=DEVICE)
424         ))
425
426     print('Total Time: %.2f min' % ((time.time() - start_time)/60))
427
428     return minibatch_cost, epoch_cost
429
430 minibatch_cost, epoch_cost = train(model, train_loader, test_loader)
431
432
433 plt.plot(range(len(minibatch_cost)), minibatch_cost)
434 plt.ylabel('Cross Entropy')
435 plt.xlabel('Minibatch')
436 plt.show()
437
438 plt.plot(range(len(epoch_cost)), epoch_cost)
439 plt.ylabel('Cross Entropy')
440 plt.xlabel('Epoch')
441 plt.show()

```

D Appendix: ResNet implementation with Python

```

1  # In[1]:
2
3
4  # pandas: plot the accuracy results, read files, save the data, etc
5  import pandas as pd
6
7  # os: make sure the code can be executed for all operating systems
8  import os
9
10 # random: draw random samples
11 import random
12
13 # imread: read and display images
14 from skimage.io import imread
15
16 # train_test_split: create validation set
17 from sklearn.model_selection import train_test_split
18
19 # pyplot: plot
20 import matplotlib.pyplot as plt
21
22 # Image: process images such as resize/clip the images
23 from PIL import Image
24
25 # shutil: remove an entire directory
26 import shutil
27
28 # time: benchmarking
29 import time
30
31 # pytorch libraries and modules

```

```

32 import torch
33
34 # numpy: process the data
35 import numpy as np
36
37 from torch.autograd import Variable
38 from torch.nn import Linear, ReLU, CrossEntropyLoss, Sequential, Conv2d, MaxPool2d, Module,
    Softmax, BatchNorm2d, Dropout
39 from torch.utils.data import Dataset
40 from torchvision import transforms
41 from torch.utils.data import DataLoader
42 import torch.nn.functional as F
43 from torch.optim import Adam, SGD
44
45
46 # In[2]:
47
48
49 #-----
50 ### SETTINGS
51 #-----
52
53 # Hyperparameters
54 RANDOM_SEED = 1
55 LEARNING_RATE = 0.05
56 BATCH_SIZE = 8
57 NUM_EPOCHS = 20
58
59 # Architecture
60 NUM_FEATURES = 32*32
61 NUM_CLASSES = 2
62 DIMENSION = (115, 61)
63
64 # Other
65 DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
66 IMG_PATH = os.path.join("dataset", "CT")
67 NUM_WORKERS = 0
68
69
70 # # Loading the data
71
72 # In[3]:
73
74
75 class CTDataset(Dataset):
76     """Custom Dataset for loading CT images"""
77
78     def __init__(self, csv_path, img_dir, transform=None):
79
80         df = pd.read_csv(csv_path)
81         self.img_dir = [img_dir] * len(df['name'].values)
82         self.img_names = df['name'].values
83         self.y = df['covid'].values
84         self.transform = transform
85
86     def __getitem__(self, index):
87         img = Image.open(os.path.join(self.img_dir[index],
88                                     self.img_names[index]))
89         #         img = imread(os.path.join(self.img_dir[index],
90         #                                     self.img_names[index]), as_gray=True).astype('
91         float32')
92         if self.transform is not None:
93             img = self.transform(img)

```

```

94         label = self.y[index]
95         return img, label
96
97     def __len__(self):
98         return self.y.shape[0]
99
100     # concatenate two datasets
101     def __add__(self, newDataset):
102         self.img_names = np.concatenate((self.img_names, newDataset.img_names))
103         self.img_dir += newDataset.img_dir
104         self.y = np.concatenate((self.y, newDataset.y))
105         return self
106
107
108 # In[4]:
109
110
111 # Note that transforms.ToTensor()
112 # already divides pixels by 255. internally
113
114 custom_transform = transforms.Compose([#transforms.Lambda(lambda x: x/255.),
115                                       transforms.ToTensor()])
116
117
118 train_dataset = CTDataset(csv_path='train.csv',
119                           img_dir=os.path.join("dataset", "CT", "train"),
120                           transform=custom_transform)
121
122 valid_dataset = CTDataset(csv_path='validation.csv',
123                            img_dir=os.path.join("dataset", "CT", "validation"),
124                            transform=custom_transform)
125
126 # this method don't need validation dataset
127
128 train_loader = DataLoader(dataset=train_dataset+valid_dataset,
129                           batch_size=BATCH_SIZE,
130                           shuffle=True,
131                           num_workers=NUM_WORKERS)
132
133
134 valid_loader = DataLoader(dataset=valid_dataset,
135                           batch_size=BATCH_SIZE,
136                           shuffle=False,
137                           num_workers=NUM_WORKERS)
138
139 test_dataset = CTDataset(csv_path='test.csv',
140                           img_dir=os.path.join("dataset", "CT", "test"),
141                           transform=custom_transform)
142
143 test_loader = DataLoader(dataset=test_dataset,
144                          batch_size=BATCH_SIZE,
145                          shuffle=False,
146                          num_workers=NUM_WORKERS)
147
148
149 # In[5]:
150
151
152 # Checking the dataset
153 for images, labels in train_loader:
154     print('Image batch dimensions:', images.shape)
155     print('Image label dimensions:', labels.shape)
156     break
157

```

```

158 # Checking the dataset
159 for images, labels in train_loader:
160     print('Image batch dimensions:', images.shape)
161     print('Image label dimensions:', labels.shape)
162     break
163
164
165 # In[6]:
166
167
168 # This cell just checks if the dataset can be loaded correctly.
169
170 torch.manual_seed(0)
171
172 num_epochs = 2
173 for epoch in range(num_epochs):
174
175     for batch_idx, (x, y) in enumerate(train_loader):
176
177         print('Epoch:', epoch+1, end='')
178         print(' | Batch index:', batch_idx, end='')
179         print(' | Batch size:', y.size()[0])
180
181         x = x.to(DEVICE)
182         y = y.to(DEVICE)
183
184         print('break minibatch for-loop')
185         break
186
187
188 # # ResNet
189
190 # In[7]:
191
192
193 def compute_epoch_loss(model, data_loader):
194     model.eval()
195     curr_loss, num_examples = 0., 0
196     with torch.no_grad():
197         for features, targets in data_loader:
198             features = features.to(DEVICE)
199             targets = targets.to(DEVICE)
200             logits = model(features)
201             loss = F.cross_entropy(logits, targets, reduction='sum')
202             num_examples += targets.size(0)
203             curr_loss += loss
204
205     curr_loss = curr_loss / num_examples
206     return curr_loss
207
208 def compute_accuracy(model, data_loader):
209     model.eval()
210     correct_pred, num_examples = 0, 0
211     for i, (features, targets) in enumerate(data_loader):
212
213         features = features.to(DEVICE)
214         targets = targets.to(DEVICE)
215
216         logits = model(features)
217         _, predicted_labels = torch.max(logits, 1)
218         num_examples += targets.size(0)
219         correct_pred += (predicted_labels == targets).sum()
220     return correct_pred.float()/num_examples * 100
221

```

```

222
223
224 # In[8]:
225
226
227 model = torch.hub.load('pytorch/vision:v0.6.0', 'resnet18', pretrained=True)
228 # or any of these variants
229 # model = torch.hub.load('pytorch/vision:v0.6.0', 'resnet34', pretrained=True)
230 # model = torch.hub.load('pytorch/vision:v0.6.0', 'resnet50', pretrained=True)
231 # model = torch.hub.load('pytorch/vision:v0.6.0', 'resnet101', pretrained=True)
232 # model = torch.hub.load('pytorch/vision:v0.6.0', 'resnet152', pretrained=True)
233 model.eval()
234
235
236 # In[9]:
237
238
239 # keep the pretrained layers, don't update them
240 for parameter in model.parameters():
241     parameter.requires_grad = False
242
243
244 # In[10]:
245
246
247 model.conv1 = torch.nn.Conv2d(1,64,kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=
    False)
248 model.fc = torch.nn.Sequential(
249     torch.nn.Linear(in_features=512, out_features=100, bias=True),
250     torch.nn.ReLU(inplace=True),
251     torch.nn.Dropout(p=0.5, inplace=False),
252     torch.nn.Linear(in_features=100, out_features= NUM_CLASSES, bias=True))
253
254
255 # In[15]:
256
257
258 # instances where you want to save and load your neural networks across different devices.
259 model = model.to(DEVICE)
260
261 # torch.optim is a package implementing various optimization algorithms.
262 optimizer = torch.optim.Adam([
263     {'params': model.conv1.parameters()},
264     {'params': model.fc.parameters()}
265 ])
266
267
268 # In[16]:
269
270
271 def train(model, train_loader, test_loader, NUM_EPOCHS):
272
273     minibatch_cost, epoch_cost = [], []
274     start_time = time.time()
275     for epoch in range(NUM_EPOCHS):
276
277         model.train()
278         for batch_idx, (features, targets) in enumerate(train_loader):
279
280             features = features.to(DEVICE)
281             targets = targets.to(DEVICE)
282
283             ### FORWARD AND BACK PROP
284             logits = model(features)

```



```

285         cost = F.cross_entropy(logits, targets)
286         optimizer.zero_grad()
287
288         cost.backward()
289         minibatch_cost.append(cost)
290
291         ### UPDATE MODEL PARAMETERS
292         optimizer.step()
293
294         ### LOGGING
295         if not batch_idx % 150:
296             print('Epoch: %03d/%03d | Batch %04d/%04d | Cost: %.4f'
297                   %(epoch+1, NUM_EPOCHS, batch_idx,
298                     len(train_loader), cost))
299
300
301         with torch.set_grad_enabled(False): # save memory during inference
302             print('Epoch: %03d/%03d | Train: %.3f%%' % (
303                 epoch+1, NUM_EPOCHS,
304                 compute_accuracy(model, train_loader)))
305
306         cost = compute_epoch_loss(model, train_loader)
307         epoch_cost.append(cost)
308
309         print('Time elapsed: %.2f min' % ((time.time() - start_time)/60))
310
311     print('Total Training Time: %.2f min' % ((time.time() - start_time)/60))
312
313
314     with torch.set_grad_enabled(False): # save memory during inference
315         print('Test accuracy: %.2f%%' % (compute_accuracy(model, test_loader)))
316
317     print('Total Time: %.2f min' % ((time.time() - start_time)/60))
318
319     return minibatch_cost, epoch_cost
320
321
322 # In[17]:
323
324
325 minibatch_cost, epoch_cost = train(model, train_loader, test_loader, NUM_EPOCHS = NUM_EPOCHS
326 )
327
328 plt.plot(range(len(minibatch_cost)), minibatch_cost)
329 plt.ylabel('Cross Entropy')
330 plt.xlabel('Minibatch')
331 plt.show()
332
333 plt.plot(range(len(epoch_cost)), epoch_cost)
334 plt.ylabel('Cross Entropy')
335 plt.xlabel('Epoch')
336 plt.show()

```

E Appendix: SVM implementation with Python

```

1 # Path: list the files in the directory
2 from pathlib import Path
3
4 # os: make sure code can be executed in all operating systems
5 import os
6

```

```

7 # time: benchmarking
8 import time
9
10 # plt: draw pictures
11 import matplotlib.pyplot as plt
12
13 # svm: svm classifier
14 from sklearn import svm, metrics
15
16 # Bunch: container object for datasets
17 from sklearn.utils import Bunch
18
19 # numpy: process the image arrays
20 import numpy as np
21
22 # pandas: process the image arrays
23 import pandas as pd
24
25 # GridSearchCV: Exhaustive search over specified parameter values for an estimator
26 # train_test_split: split the full datasets into train and test dataset, respectively
27 from sklearn.model_selection import GridSearchCV, train_test_split
28
29 # imread: read the images
30 from skimage.io import imread
31
32 # resize: resize the images
33 from skimage.transform import resize
34
35 # plot_decision_regions: visualize the SVM hyperplane
36 from mlxtend.plotting import plot_decision_regions
37
38 DIMENSION = (115, 61)
39 IMG_PATH = os.path.join("dataset", "CT")
40
41 def load_image_files(container_path, dimension=DIMENSION):
42     """
43     Load image files with categories as subfolder names
44     which performs like scikit-learn sample dataset
45
46     Parameters
47     -----
48     container_path : string or unicode
49         Path to the main folder holding one subfolder per category
50     dimension : tuple
51         size to which image are adjusted to
52
53     Returns
54     -----
55     Bunch
56     """
57
58     image_dir = Path(container_path)
59     folders = [directory for directory in image_dir.iterdir() if (directory.is_dir() and "
COVID" in directory.name)]
60     categories = [fo.name for fo in folders]
61
62     descr = "A image classification dataset"
63     images = []
64     flat_data = []
65     target = []
66     for i, direc in enumerate(folders):
67         for file in direc.iterdir():
68             if len(file.name.split(".")) == 1 or file.name[0] == '.':
69                 continue

```

```

70         img = imread(file, as_gray = True)
71         img_resized = resize(img, dimension, anti_aliasing=True, mode='reflect')
72         flat_data.append(img_resized.flatten())
73         images.append(img_resized)
74         target.append(i)
75     flat_data = np.array(flat_data)
76     target = np.array(target)
77     images = np.array(images)
78
79     return Bunch(data=flat_data,
80                 target=target,
81                 target_names=categories,
82                 images=images,
83                 DESCR=descr)
84
85 image_dataset = load_image_files(IMG_PATH)
86
87 X_train, X_test, y_train, y_test = train_test_split(
88     image_dataset.data, image_dataset.target, test_size=0.1, random_state=109)
89
90 param_grid = [
91     {'C': [1, 10, 100, 1000], 'kernel': ['linear']},
92     {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001], 'kernel': ['rbf']},
93 ]
94 svc = svm.SVC()
95
96 start = time.time()
97 clf1 = GridSearchCV(svc, param_grid, cv = 3)
98 clf1.fit(X_train, y_train)
99 y_pred = clf1.predict(X_test)
100 print("Classification report for - \n{}:\n{}\n".format(
101     clf1, metrics.classification_report(y_test, y_pred)))
102 print("Time elapsed: %.2f min" % ((time.time()-start)/60))
103
104 clf1.best_params_
105
106 param_grid = [
107     {'C': [10, 50], 'gamma': [0.002, 0.001, 0.0005], 'kernel': ['rbf']},
108 ]
109 svc = svm.SVC()
110
111 start = time.time()
112 clf2 = GridSearchCV(svc, param_grid, cv = 3)
113 clf2.fit(X_train, y_train)
114 y_pred = clf2.predict(X_test)
115 print("Classification report for - \n{}:\n{}\n".format(
116     clf2, metrics.classification_report(y_test, y_pred)))
117 print("Time elapsed: %.2f min" % ((time.time()-start)/60))
118
119 clf2.best_params_
120
121 # train accuracy
122 y_pred = clf2.predict(X_train)
123 print("Classification report for - \n{}:\n{}\n".format(
124     clf1, metrics.classification_report(y_train, y_pred)))

```